

Chapter 9: Memory Management

- Background
- General Memory Management
- Compiling, Linking, and Loading
- Swapping
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table
- Segmentation

Background

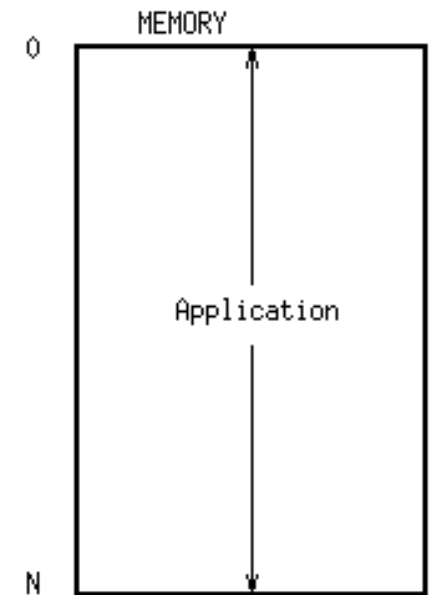
- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are the only storage CPU can access directly
- Access time differs:
 - Register access in one CPU clock (or less)
 - Main memory can take many cycles
- Cache sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

Memory Management

- Several questions we can ask about our O/S:
 - What restrictions does the operating system and hardware place on user (process) access to memory?
 - What facilities are provided by the operating system and hardware for memory access by processes?
 - What protection is given by the memory management scheme?
- Many different memory management schemes...

Bare Machine (No Memory Management)

- The application has complete control over all the memory in the machine
- No separate operating system is present
- There is no protection of any memory location
 - the application process can write to any memory location

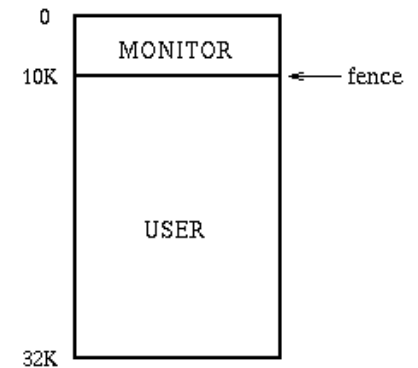


Resident Monitor

➤ Memory is divided into 2 parts:

➤ the monitor (or operating system)

➤ the user (or application)



➤ The monitor is generally placed at the top or the bottom of memory

➤ We must have a way of protecting the operating system code from corruption

➤ application may either maliciously or accidentally write into operating system space

Resident Monitor

- Memory protection can be implemented in hardware using a fence:
 - A memory access by the user program on the OS side of the fence generates a hardware trap (kind of like a segmentation fault)
 - Depending on the hardware, this can slow down the execution as each memory reference must be checked.
 - This requires 2 modes of execution
 - **User mode:** hardware checks all memory references
 - **Supervisor mode:** hardware checks disabled

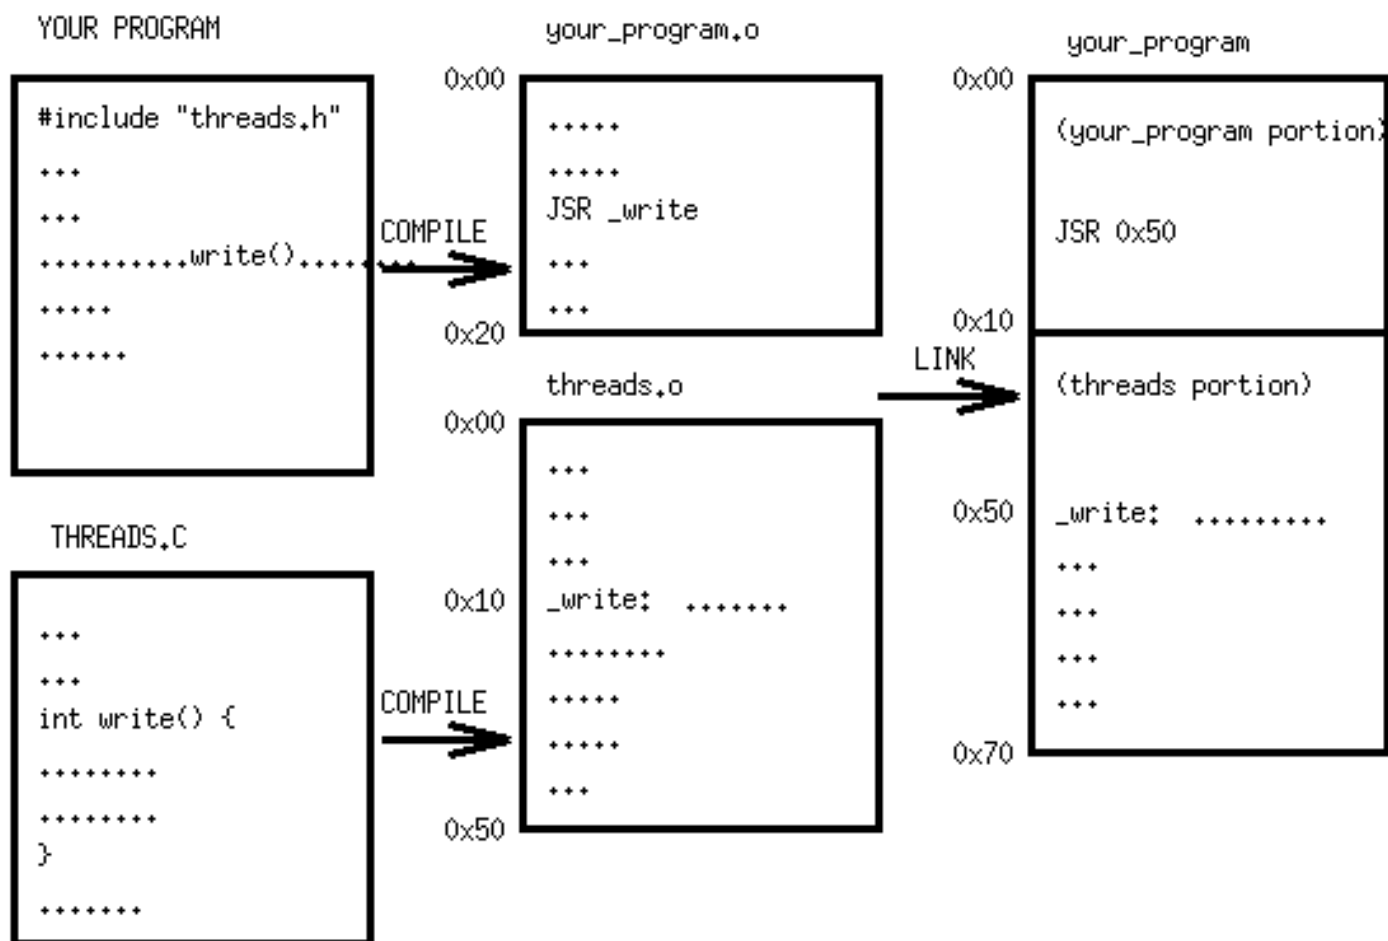
Problems with Resident Monitors

- Early monitors: the fence had to be built into the hardware
 - The fence remained constant
 - Constrained O/S size: Either some memory was wasted, or part of the O/S was unprotected
- Later: the fence value was placed in a register
 - The fence register could accommodate the size of the O/S
- Fence location change \Rightarrow problems:
 - means re-linking programs to set addresses

Compiling, Linking, and Loading

- compiling: **source** → **object file**
- linking: **object files** → **load module**
 - resolve addresses of references (internal and external)
 - if the load module will be loaded at the same address each time, the linker can generate absolute addresses
 - otherwise, relocatable addresses/objects
- loading: **load module** → **main memory**

Linking



Static Relocation

- The linker always generates a module with addresses assuming the module will be loaded at address zero
- The loader scans the module and adjusts every address by adding the actual load location to each address
- The loader must be "smart" enough to be able to interpret the instructions in the load module - at least enough to identify the addressing modes of the instructions it finds
- Once a module has begun execution it cannot be relocated

Dynamic Relocation

- In this scheme, the linker still generates addresses assuming a start position of zero
- MMU updates each memory reference by adding the actual load location to the memory address as the program runs
- An application never sees a *physical memory address* - only the *logical memory addresses* generated by the linker
 - Conversion between physical and logical addresses is done by the Memory Management Unit (MMU)
- The loader need not be "smart" - it simply copies the load module into its load location (simpler, faster)