

# General Dynamic Storage Allocation Problem

## ➤ Problem:

- allocating contiguous blocks of storage from an initially free, single, large block

## ➤ Goals:


- efficiency of allocation and freeing
- minimization of fragmentation

# General Dynamic Storage Allocation Problem

- **First Fit:** allocate the first block on the list that is big enough to satisfy the request.
  - this method is fast
- **Best Fit:** allocate the smallest free block bigger than the block size that is needed.
  - Tries to reduce fragmentation by minimizing left over space
- **Worst Fit:** allocate the largest block available.
  - Tries to reduce fragmentation by leaving the most usable leftover piece

# General Dynamic Storage Allocation Problem

➤ What if a request is made for 19995 blocks of storage, and there is a 20000 block space free?



➤ The overhead to keep track of the wasted five bytes is more than it is worth (more than 5 bytes needed to track a memory block)

➤ **allocate all 20000 blocks**

➤ **produces a very small amount of internal fragmentation**

➤ There are many algorithms for the general storage allocation problem


➤ Many trade off one of allocation/freeing time and fragmentation for the other

➤ Some of these algorithms are, or are close to  $O(1)$  in both allocation and freeing (at the expense of fragmentation)

# Compaction

➤ With MVT, external fragmentation can get very severe.

➤ e.g a small free block on both sides of every used block



➤ Compaction:

➤ shuffle all used blocks of memory into one, contiguous, large block

# Compaction

- External fragmentation eliminated by combining all fragments into one large block
  - may be big enough to run a program while each of the smaller ones were not
- Can compaction always be done?  
What is necessary to make it possible?
  - we need dynamic relocation
    - logical addresses must be bound to physical addresses at run time so we can move the program in memory, once execution has begun
  - We need only move the program and change the base register(s)


# Compaction

## ➤ Alternatives:

- move job 4 to the slot above job 3 (copy only 400K)
- move job 3 to below job four, leaving a large, 900K hole in the middle (copy only 200K)
- move both jobs 3 and 4 up

## ➤ Swapping / Compaction can be used together

- When the jobs are swapped back in, they can be placed together to leave the largest hole.



# Paging

- Fragmentation is a big problem with MVT
  - The memory occupied by a program must be contiguous
- We could greatly reduce the cost of dealing with fragmentation if the memory didn't need to be contiguous
- Paging allows a logical block to be scattered throughout physical memory
  - any available chunk of memory is usable even if it isn't large enough to hold the entire object

# Paging

- Requires hardware support
- The load module is broken into a series of fixed-size pages
  - size determined by hardware (2k and 4k are common)
  - all pages but the last page are the same sizes
    - e.g. a 15K program might be broken down into three 4k pages and one 3K page
- Main memory divided into frames
  - frame size = page size
  - one page fits into one frame exactly
  - last page for each object may be smaller, therefore occupies only part of one frame (remainder is wasted)



# Paging

➤ Each logical address consists of two parts:

➤ page number: 
$$= (\text{offset from top of module}) / (\text{page size})$$

➤ displacement within the page:  
$$= (\text{offset from top of module}) \% (\text{page size}) \text{ (modulo)}$$

# Paging

➤ How do we derive these page numbers and displacements?

➤ We use a page size of  $2^N$

➤ Displacement → least significant  $N$  bits of logical address

➤ Page number → remaining (most significant) bits

➤ E.g.      logical address



➤ Why do we do it this way?

➤ Can determine page # and displacement efficiently in hardware

# Paging

➤ **page table:** data structure for each process

➤ entries for the base address of each page in physical memory

➤ map showing where each part of the process is loaded in memory

load  
module

P0
P1
P2
P3

page  
table


0	1
1	4
2	3
3	7

main  
memory

FRAME NUMBER		
0		
1	P0	
2		
3	P2	<—
4	P1	
5		
6		
7	P3	<—

# Paging

- physical address calculation:
  - use page number as an index into the page table
    - derives the base address of the frame
  - add the displacement to derive the physical address



# Paging

- When a job is to be loaded:
  - We see how many frames are required for the job
  - If enough frames are available, we load the job into the free frames, and update the page table for this job
- What about fragmentation with this paging scheme?  
Can we have:
  - External fragmentation? **yes, but very little**
  - Internal fragmentation? **yes, but only for last page of each process**
- The page table for each job is kept in its PCB
  - There's also a hardware register that has a pointer to the page table for the currently executing process

# Page Table Implementation

- Where do we store the page table?
  - hardware registers – theres not enough X
  - use memory + cache called Translation Lookaside Buffer
- Average memory access time:
  - cache lookup time = 50 ns
  - main memory access time = 750 ns
  - cache hit ratio is 80%
  - avg memory access time =  
 $0.8 * (50 + 750) + 0.2 * (50 + 750 + 750) = 950 \text{ ns} \sim 26\%$
  - slowdown compared to non-paging lookup

# Page Table Implementation

➤ What if we increase the cache size?

E.g hit ratio of 95%

avg memory access time =

$$(0.95) * (50 + 750) + 0.05 * (50 + 750 + 750) = 837.5 \text{ ns}$$

~ 11% slowdown

➤ Paging allows code (pages) to be shared

➤ E.g. editor can be loaded by the first user to load it, and subsequent users page table will point to the previously loaded code

➤ code must be "reentrant"

➤ Paging separates logical (contiguous) memory and physical memory